

resnet18_changeplot

June 13, 2020

```
[2]: import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
import torchvision
import torchvision.models as models

import os
import torch
import librosa

import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline

from datasets import AudioDataset
```

0.0.1 package path

```
[3]: # path to urban sound 8k
data_root = "data/UrbanSound8K/"
# path to label
label_path = "data/UrbanSound8K/metadata/UrbanSound8K.csv"
```

0.0.2 learning rate

```
[4]: def lr_schedule(epoch, init_lr):
    if epoch < 20:
        return init_lr
    elif epoch >= 20 and epoch < 40:
        return init_lr/10
    elif epoch >= 40 and epoch < 80:
        return init_lr/100
    else:
        return init_lr/1000
```

0.0.3 model construction

```
[5]: class ResidualBlock(nn.Module):
    def __init__(self, inchannel, outchannel, stride, shortcut=None):
        super(ResidualBlock, self).__init__()
        self.basic = nn.Sequential(
            nn.Conv2d(inchannel, outchannel, 3, stride, 1, bias=False), # stride
            nn.BatchNorm2d(outchannel), #
            nn.ReLU(inplace=True), #
            nn.Conv2d(outchannel, outchannel, 3, 1, 1, bias=False), # feature_
            ↪map
            nn.BatchNorm2d(outchannel),
        )
        self.shortcut = shortcut

    def forward(self, x):
        out = self.basic(x)
        residual = x if self.shortcut is None else self.shortcut(x) #
        out += residual
        return nn.ReLU(inplace=True)(out) #

#ResNet
class ResNet(nn.Module):
    def __init__(self):
        super(ResNet, self).__init__()
        self.pre = nn.Sequential(
            nn.Conv2d(3, 64, 7, 2, 3, bias=False),
            nn.BatchNorm2d(64),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(3, 2, 1),
        ) #
        self.body = self.makelayers([3, 4, 6, 3]) #
        self.classifier = nn.Linear(512, 10) #

    def makelayers(self, blocklist): #
        self.layers = []
        for index, blocknum in enumerate(blocklist):
            if index != 0:
                shortcut = nn.Sequential(
                    nn.Conv2d(64*2**(index-1), 64*2**index, 1, 2, bias=False),
                    nn.BatchNorm2d(64*2**index)
                ) #
                self.layers.
                ↪append(ResidualBlock(64*2**(index-1), 64*2**index, 2, shortcut)) #
                for i in range(0 if index==0 else 1, blocknum):
                    self.layers.append(ResidualBlock(64*2**index, 64*2**index, 1))
        return nn.Sequential(*self.layers)
```

```

def forward(self,x):
    x = self.pre(x)
    x = self.body(x)
    x = nn.AvgPool2d(7)(x) #kernel_size 7      feature_
↪map 7*7 224->112->56->28->14->7
    x = x.view(x.size(0),-1)
    x = self.classifier(x)
    return x

```

```

[6]: net = ResNet()
     print(net)

```

```

ResNet(
  (pre): Sequential(
    (0): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3),
bias=False)
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (2): ReLU(inplace=True)
    (3): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1,
ceil_mode=False)
  )
  (body): Sequential(
    (0): ResidualBlock(
      (basic): Sequential(
        (0): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
        (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU(inplace=True)
        (3): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
        (4): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
    (1): ResidualBlock(
      (basic): Sequential(
        (0): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
        (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU(inplace=True)
        (3): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
        (4): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

```

```

    )
)
(2): ResidualBlock(
  (basic): Sequential(
    (0): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (2): ReLU(inplace=True)
    (3): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
    (4): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  )
)
(3): ResidualBlock(
  (basic): Sequential(
    (0): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1),
bias=False)
    (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (2): ReLU(inplace=True)
    (3): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
    (4): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  )
  (shortcut): Sequential(
    (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
    (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  )
)
(4): ResidualBlock(
  (basic): Sequential(
    (0): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
    (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (2): ReLU(inplace=True)
    (3): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
    (4): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  )
)
(5): ResidualBlock(
  (basic): Sequential(

```

```

        (0): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
        (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU(inplace=True)
        (3): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
        (4): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
)
(6): ResidualBlock(
  (basic): Sequential(
    (0): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
    (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (2): ReLU(inplace=True)
    (3): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
    (4): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  )
)
(7): ResidualBlock(
  (basic): Sequential(
    (0): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1),
bias=False)
    (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (2): ReLU(inplace=True)
    (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
    (4): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  )
  (shortcut): Sequential(
    (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
    (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  )
)
(8): ResidualBlock(
  (basic): Sequential(
    (0): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
    (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

```

```

        (2): ReLU(inplace=True)
        (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
        (4): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
)
    (9): ResidualBlock(
        (basic): Sequential(
            (0): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
            (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): ReLU(inplace=True)
            (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
            (4): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
    )
    (10): ResidualBlock(
        (basic): Sequential(
            (0): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
            (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): ReLU(inplace=True)
            (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
            (4): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
    )
    (11): ResidualBlock(
        (basic): Sequential(
            (0): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
            (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): ReLU(inplace=True)
            (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
            (4): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
    )
    (12): ResidualBlock(
        (basic): Sequential(

```

```

        (0): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
        (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU(inplace=True)
        (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
        (4): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
)
(13): ResidualBlock(
  (basic): Sequential(
    (0): Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1),
bias=False)
    (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (2): ReLU(inplace=True)
    (3): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
    (4): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  )
  (shortcut): Sequential(
    (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
    (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  )
)
(14): ResidualBlock(
  (basic): Sequential(
    (0): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
    (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (2): ReLU(inplace=True)
    (3): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
    (4): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  )
)
(15): ResidualBlock(
  (basic): Sequential(
    (0): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
    (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

```

```

        (2): ReLU(inplace=True)
        (3): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
        (4): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
)
(classifier): Linear(in_features=512, out_features=10, bias=True)
)

```

```
[7]: print(models.resnet34(num_classes=10))
```

```

ResNet(
  (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3),
bias=False)
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (relu): ReLU(inplace=True)
  (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1,
ceil_mode=False)
  (layer1): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (1): BasicBlock(
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (2): BasicBlock(
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,

```



```

track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
    (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (layer2): Sequential(
    (0): BasicBlock(
        (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (downsample): Sequential(
            (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
            (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
    )
    (1): BasicBlock(
        (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (2): BasicBlock(
        (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (3): BasicBlock(

```

```

        (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
)
(layer3): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  )
  (2): BasicBlock(
    (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,

```

```

track_running_stats=True)
    )
    (3): BasicBlock(
      (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (4): BasicBlock(
      (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (5): BasicBlock(
      (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (layer4): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
      (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
      (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (downsample): Sequential(

```

```

        (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
        (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
)
(1): BasicBlock(
  (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
  (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (relu): ReLU(inplace=True)
  (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
  (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
)
(2): BasicBlock(
  (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
  (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (relu): ReLU(inplace=True)
  (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
  (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
)
)
(avgpool): AdaptiveAvgPool2d(output_size=(1, 1))
(fc): Linear(in_features=512, out_features=10, bias=True)
)

```

0.0.4 training fuction

```

[14]: def train(MAX_EPOCH = 100):

#     eval_interval = 5

# initialize dataset (feature can be "mfcc", "spec", "mel_raw")
audio_dataset = AudioDataset(5, DataRoot=data_root, LabelPath=label_path,
↪feature="mfcc", mode="train")

# define resnet model
cnn_model = models.resnet18(num_classes=10)
#cnn_model =ResNet()
# to gpu

```

```

cnn_model = cnn_model.cuda()
#print(cnn_model)
# initialize dataloader
data_loader = torch.utils.data.DataLoader(audio_dataset, batch_size=32,
→shuffle=True, num_workers=1)

# loss function
loss_fn = nn.CrossEntropyLoss().cuda()

# lr
learning_rate = 1e-3

# initialize optimizer
optimizer = torch.optim.Adam(cnn_model.parameters(), lr=learning_rate)

# initialize logger
train_acc = []
test_acc = []

train_loss = []
test_loss = []

for epoch in range(MAX_EPOCH):
    # iterate through dataset
    for param_group in optimizer.param_groups:
        param_group['lr'] = lr_schedule(epoch, learning_rate)

    # initialize epoch stat
    correct_num = 0
    total_num = 0
    loss_sum = 0

    for idx, data in enumerate(data_loader):
        #print(idx)
        train_data, labels = data

        #train_data = train_data.type(torch.float32)/255

        # data to gpu
        train_data = train_data.cuda()
        labels = labels.cuda()

        prob = cnn_model(train_data)
        loss = loss_fn(prob, labels)

        output = prob.argmax(1)

```

```

        loss_sum += loss.item()*float(labels.shape[0])
        correct_num += (output==labels).sum().double()
        total_num += float(labels.shape[0])

    optimizer.zero_grad()

    loss.backward()

    optimizer.step()

    train_acc.append(correct_num/total_num)
    train_loss.append(loss_sum/total_num)

    print("epoch: {} acc: {:.4} avg loss: {:.4f}".format(epoch, correct_num/
→total_num, loss_sum/total_num))

#         if epoch%5 == 4:

            a,b = test(cnn_model)
            test_acc.append(a)
            test_loss.append(b)

plt.figure()
plt.plot(np.arange(MAX_EPOCH), train_acc)
#     plt.plot(np.arange(eval_interval-1, MAX_EPOCH, eval_interval), test_acc)
plt.plot(np.arange(MAX_EPOCH), test_acc)
plt.title("accuracy")
plt.legend(["train", "val"])

plt.figure()
plt.plot(np.arange(MAX_EPOCH), train_loss)
#     plt.plot(np.arange(eval_interval-1, MAX_EPOCH, eval_interval), test_loss)
plt.plot(np.arange(MAX_EPOCH), test_loss)
plt.title("loss")
plt.legend(["train", "val"])

```

0.0.5 Test function

```

[15]: def test(model, ):
        test_dataset = AudioDataset(5, DataRoot=data_root, LabelPath=label_path,
→feature="mfcc", mode="test")
        test_loader = torch.utils.data.DataLoader(test_dataset, batch_size=32,
→shuffle=True, num_workers=1)

        model.eval()

```

```

loss_fn = nn.CrossEntropyLoss().cuda()

correct_num = 0
total_num = 0
loss_sum = 0

for idx, data in enumerate(test_loader):
    test_data, labels = data
    test_data = test_data.cuda()
    labels = labels.cuda()

    test_data = test_data.cuda()

    prob = model(test_data)
    loss = loss_fn(prob, labels)

    output = prob.argmax(1)

    correct_num += (output==labels).sum().double()
    loss_sum += loss.item()*float(labels.shape[0])
    total_num += float(labels.shape[0])

model.train()

print("##Testing## epoch acc: {:.4}".format(correct_num/total_num))
return correct_num/total_num, loss_sum/total_num

```

```
[16]: train(70)
```

```

verify mfcc feature success
epoch: 0 acc: 0.2199 avg loss: 2.1515
verify mfcc feature success
##Testing## epoch acc: 0.2201
epoch: 1 acc: 0.2851 avg loss: 1.9610
verify mfcc feature success
##Testing## epoch acc: 0.1763
epoch: 2 acc: 0.3638 avg loss: 1.7972
verify mfcc feature success
##Testing## epoch acc: 0.3387
epoch: 3 acc: 0.4433 avg loss: 1.5925
verify mfcc feature success
##Testing## epoch acc: 0.3782
epoch: 4 acc: 0.5249 avg loss: 1.3813
verify mfcc feature success
##Testing## epoch acc: 0.3985
epoch: 5 acc: 0.6026 avg loss: 1.1548

```

```
verify mfcc feature success
##Testing## epoch acc: 0.5214
epoch: 6 acc: 0.6528 avg loss: 1.0095
verify mfcc feature success
##Testing## epoch acc: 0.4786
epoch: 7 acc: 0.6997 avg loss: 0.8901
verify mfcc feature success
##Testing## epoch acc: 0.5256
epoch: 8 acc: 0.7376 avg loss: 0.7779
verify mfcc feature success
##Testing## epoch acc: 0.5737
epoch: 9 acc: 0.7594 avg loss: 0.6997
verify mfcc feature success
##Testing## epoch acc: 0.5556
epoch: 10 acc: 0.7748 avg loss: 0.6646
verify mfcc feature success
##Testing## epoch acc: 0.5224
epoch: 11 acc: 0.7986 avg loss: 0.5944
verify mfcc feature success
##Testing## epoch acc: 0.5855
epoch: 12 acc: 0.819 avg loss: 0.5373
verify mfcc feature success
##Testing## epoch acc: 0.6015
epoch: 13 acc: 0.8284 avg loss: 0.5023
verify mfcc feature success
##Testing## epoch acc: 0.5299
epoch: 14 acc: 0.8406 avg loss: 0.4581
verify mfcc feature success
##Testing## epoch acc: 0.6453
epoch: 15 acc: 0.8556 avg loss: 0.4220
verify mfcc feature success
##Testing## epoch acc: 0.5534
epoch: 16 acc: 0.8593 avg loss: 0.4047
verify mfcc feature success
##Testing## epoch acc: 0.5791
epoch: 17 acc: 0.8763 avg loss: 0.3589
verify mfcc feature success
##Testing## epoch acc: 0.5192
epoch: 18 acc: 0.8837 avg loss: 0.3414
verify mfcc feature success
##Testing## epoch acc: 0.5363
epoch: 19 acc: 0.8849 avg loss: 0.3244
verify mfcc feature success
##Testing## epoch acc: 0.5449
epoch: 20 acc: 0.9312 avg loss: 0.2049
verify mfcc feature success
##Testing## epoch acc: 0.6015
epoch: 21 acc: 0.945 avg loss: 0.1666
```



```
verify mfcc feature success
##Testing## epoch acc: 0.6197
epoch: 22 acc: 0.9541 avg loss: 0.1447
verify mfcc feature success
##Testing## epoch acc: 0.6036
epoch: 23 acc: 0.9531 avg loss: 0.1425
verify mfcc feature success
##Testing## epoch acc: 0.6004
epoch: 24 acc: 0.9574 avg loss: 0.1335
verify mfcc feature success
##Testing## epoch acc: 0.6079
epoch: 25 acc: 0.9616 avg loss: 0.1199
verify mfcc feature success
##Testing## epoch acc: 0.6004
epoch: 26 acc: 0.9616 avg loss: 0.1234
verify mfcc feature success
##Testing## epoch acc: 0.6079
epoch: 27 acc: 0.961 avg loss: 0.1166
verify mfcc feature success
##Testing## epoch acc: 0.5951
epoch: 28 acc: 0.9649 avg loss: 0.1081
verify mfcc feature success
##Testing## epoch acc: 0.5887
epoch: 29 acc: 0.9683 avg loss: 0.1015
verify mfcc feature success
##Testing## epoch acc: 0.609
epoch: 30 acc: 0.9687 avg loss: 0.0975
verify mfcc feature success
##Testing## epoch acc: 0.6132
epoch: 31 acc: 0.9701 avg loss: 0.0944
verify mfcc feature success
##Testing## epoch acc: 0.6026
epoch: 32 acc: 0.971 avg loss: 0.0947
verify mfcc feature success
##Testing## epoch acc: 0.6132
epoch: 33 acc: 0.9709 avg loss: 0.0939
verify mfcc feature success
##Testing## epoch acc: 0.5908
epoch: 34 acc: 0.9772 avg loss: 0.0764
verify mfcc feature success
##Testing## epoch acc: 0.5908
epoch: 35 acc: 0.9774 avg loss: 0.0743
verify mfcc feature success
##Testing## epoch acc: 0.6175
epoch: 36 acc: 0.9777 avg loss: 0.0733
verify mfcc feature success
##Testing## epoch acc: 0.6058
epoch: 37 acc: 0.9792 avg loss: 0.0706
```

```
verify mfcc feature success
##Testing## epoch acc: 0.6004
epoch: 38 acc: 0.9785 avg loss: 0.0691
verify mfcc feature success
##Testing## epoch acc: 0.61
epoch: 39 acc: 0.9801 avg loss: 0.0642
verify mfcc feature success
##Testing## epoch acc: 0.6154
epoch: 40 acc: 0.9846 avg loss: 0.0543
verify mfcc feature success
##Testing## epoch acc: 0.6132
epoch: 41 acc: 0.9801 avg loss: 0.0612
verify mfcc feature success
##Testing## epoch acc: 0.6154
epoch: 42 acc: 0.986 avg loss: 0.0500
verify mfcc feature success
##Testing## epoch acc: 0.6132
epoch: 43 acc: 0.9838 avg loss: 0.0539
verify mfcc feature success
##Testing## epoch acc: 0.6004
epoch: 44 acc: 0.9863 avg loss: 0.0502
verify mfcc feature success
##Testing## epoch acc: 0.6143
epoch: 45 acc: 0.9845 avg loss: 0.0546
verify mfcc feature success
##Testing## epoch acc: 0.61
epoch: 46 acc: 0.9845 avg loss: 0.0508
verify mfcc feature success
##Testing## epoch acc: 0.6036
epoch: 47 acc: 0.9858 avg loss: 0.0491
verify mfcc feature success
##Testing## epoch acc: 0.6111
epoch: 48 acc: 0.986 avg loss: 0.0492
verify mfcc feature success
##Testing## epoch acc: 0.609
epoch: 49 acc: 0.9878 avg loss: 0.0461
verify mfcc feature success
##Testing## epoch acc: 0.6218
epoch: 50 acc: 0.9856 avg loss: 0.0503
verify mfcc feature success
##Testing## epoch acc: 0.6207
epoch: 51 acc: 0.9879 avg loss: 0.0461
verify mfcc feature success
##Testing## epoch acc: 0.6111
epoch: 52 acc: 0.9891 avg loss: 0.0438
verify mfcc feature success
##Testing## epoch acc: 0.609
epoch: 53 acc: 0.9879 avg loss: 0.0457
```

```
verify mfcc feature success
##Testing## epoch acc: 0.6154
epoch: 54 acc: 0.9878 avg loss: 0.0459
verify mfcc feature success
##Testing## epoch acc: 0.6036
epoch: 55 acc: 0.9872 avg loss: 0.0476
verify mfcc feature success
##Testing## epoch acc: 0.6036
epoch: 56 acc: 0.9876 avg loss: 0.0450
verify mfcc feature success
##Testing## epoch acc: 0.6143
epoch: 57 acc: 0.9877 avg loss: 0.0434
verify mfcc feature success
##Testing## epoch acc: 0.6229
epoch: 58 acc: 0.9883 avg loss: 0.0462
verify mfcc feature success
##Testing## epoch acc: 0.6132
epoch: 59 acc: 0.987 avg loss: 0.0449
verify mfcc feature success
##Testing## epoch acc: 0.6047
epoch: 60 acc: 0.9879 avg loss: 0.0447
verify mfcc feature success
##Testing## epoch acc: 0.5994
epoch: 61 acc: 0.9865 avg loss: 0.0458
verify mfcc feature success
##Testing## epoch acc: 0.6132
epoch: 62 acc: 0.9886 avg loss: 0.0429
verify mfcc feature success
##Testing## epoch acc: 0.6165
epoch: 63 acc: 0.9869 avg loss: 0.0476
verify mfcc feature success
##Testing## epoch acc: 0.6175
epoch: 64 acc: 0.9878 avg loss: 0.0414
verify mfcc feature success
##Testing## epoch acc: 0.6058
epoch: 65 acc: 0.9873 avg loss: 0.0427
verify mfcc feature success
##Testing## epoch acc: 0.6165
epoch: 66 acc: 0.9903 avg loss: 0.0401
verify mfcc feature success
##Testing## epoch acc: 0.6015
epoch: 67 acc: 0.9876 avg loss: 0.0431
verify mfcc feature success
##Testing## epoch acc: 0.61
epoch: 68 acc: 0.9903 avg loss: 0.0390
verify mfcc feature success
##Testing## epoch acc: 0.6143
epoch: 69 acc: 0.9906 avg loss: 0.0397
```

verify mfcc feature success
##Testing## epoch acc: 0.6079

